

The Point of Mens Rea: The Case of Willful Ignorance

Gideon Yaffe¹ 

Published online: 25 October 2016

© Springer Science+Business Media Dordrecht 2016

Abstract Under the “Willful Ignorance Principle,” a defendant is guilty of a crime requiring knowledge he lacks provided he is ignorant thanks to having earlier omitted inquiry. In this paper, I offer a novel justification of this principle through application of the theory that knowledge matters to culpability because of how the knowing action manifests the agent’s failure to grant sufficient weight to other people’s interests. I show that, under a simple formal model that supports this theory, omitting inquiry manifests precisely the same degree of disregard of others’ interests as manifested in knowingly acting criminally. Several surprising implications of this view are described, including that when the agent’s method of inquiry has a non-zero false positive rate, his omission of inquiry does not make the same contribution to his culpability as knowledge, while it does, by contrast, when the false negative rate is non-zero.

Keywords Willful ignorance · Self-deception · Mens rea · Criminal culpability · Knowledge

1 Introduction

Carmen Heredia took the bus with her mother from Tucson to Nogales, Mexico for a dentist appointment. Her two young daughters joined them. On the return trip, they borrowed her Aunt Belia’s car. Her aunt joined the four of them for the trip back into the United States, Heredia behind the wheel. The suspicion of border patrol agents was aroused by the strong smell of detergent in the car. A search turned up nearly 350 lb of marijuana stored in the trunk and wrapped in dryer sheets, apparently to hide the odor of the marijuana. Heredia appeared to be shocked to see the marijuana in the car; she seemed to have

✉ Gideon Yaffe
gideon.yaffe@yale.edu

¹ Yale Law School, New Haven, CT, USA

had no idea it was there. However, she had had a few opportunities to learn of the presence of the marijuana in the car she would be driving, and she admitted to being suspicious. Her aunt had offered an implausible explanation for the odor, and Heredia did not believe her; her mother seemed very nervous and had a surprisingly large amount of cash, despite being unemployed. But Heredia never looked in the trunk and never queried her mother about the money; she never took the steps she would have needed to take to learn that there was marijuana in the car. Hers is a sympathetic case. Her mother and her aunt were using not just Heredia, but also her children, to cloak their efforts to move drugs into the country. What could be more innocent than a car full of children and the smell of fresh laundry? And yet Heredia was convicted of *knowingly* transporting marijuana across the border, a conviction that was upheld on appeal.¹

One might have thought that if a statute says that doing something knowingly is a crime, and a particular person did not have the requisite knowledge, then the Principle of Legality, according to which there can be no punishment without a violation of the law, would preclude conviction of the person for the violation of that law. Maybe he could be convicted of violating some other law, but certainly not *that* one. But, in fact, this is not so.² The result in Heredia's case is entirely consistent with the Principle of Legality, given the presence in the law of *mens rea* substitution principles: principles that tell us that, when a particular mental state, such as knowledge, is an element of a crime, the prosecution can meet its burden by showing something else. In this case, the guilty verdict rests on the following principle:

The Willful Ignorance Principle: If a defendant was ignorant of a fact when he acted, and his ignorance was a result of his failure to inquire at an earlier time about whether the fact would be present at the time of action, then he is to be given the same legal treatment as someone who knew the fact at the time of action.

By showing that Heredia's ignorance of the fact that the car contained marijuana was due to her failure to inquire about that, the state meets its burden of establishing the element of knowledge while, at the same time, granting that Heredia was ignorant of the crucial facts.

The justification of a *mens rea* substitution principle, like the Willful Ignorance Principle, necessarily requires an account of the contribution of the statutorily identified mental state to the rationale for responding with criminal liability and punishment to those who meet the statute's definition of the crime. This is essential because it needs to be shown that, when the conditions of substitution are met, we have the same rationale for an assignment of criminal liability as we have in cases in which the statutorily specified mental state is present. To justify a *mens rea* substitution principle, then, we need to know why *mens rea* matters; we need to know *its point*. We need to know what explanatory contribution facts about a person's psychology at the time of bad behavior make to the imposition of criminal liability for that behavior. Armed with such an account, we can then see what other, distinct psychological facts make the same explanatory contribution. By identifying such facts, we will know which psychological states are appropriately substituted for those that are explicitly made necessary conditions of crime by statute.

¹ *United States v. Heredia* (483 F.3d 913).

² Both the description offered here of the apparent conflict between willful ignorance doctrine and the Principle of Legality, and the explanation for why the conflict is only apparent, follow Douglas Husak and Craig Callender (1994), "Wilful Ignorance, Knowledge and the 'Equal Culpability Thesis': A Study in the Deeper Significance of the Principle of Legality," *Wisconsin Law Review* 1: 29–68.

My goal here is to advance a general theory of the importance of mental state to criminal liability.³ In pursuit of that goal, I will offer a rationale for the Willful Ignorance Principle.

Those familiar with the law in this domain will recognize that as just formulated the principle overstates what the law actually is in any jurisdiction in the United States. No one will be treated as knowing just in virtue of the fact that he failed to inquire. Other conditions need to be met. Some jurisdictions, for instance, require that the defendant failed to inquire *so as to avoid prosecution for a knowledge crime*.⁴ Others require that the defendant failed to inquire *while recognizing a substantial risk that the relevant condition was in place*.⁵ Others, following the Model Penal Code, require only that the defendant failed to inquire *while aware of a high probability and while lacking a belief that the relevant condition was absent*.⁶ These are all efforts to capture the set of conditions under which it is justified to apply the Willful Ignorance Principle in the unqualified form in which I have here stated it.⁷

I will explain here a set of conditions under which the Willful Ignorance Principle is justified and, in so doing, I will offer a general theory of the importance of mental state to criminal responsibility. As we will see, in certain conditions, to be identified here, willfully ignorant actors are just the same as knowing actors with respect to the features thanks to which knowing actors are held criminally liable and punished. To see this, we will first need to identify the features of knowing actors that contribute to their culpability for

³ For the same view expounded in a different way, see Gideon Yaffe (2012), “Intoxication, Recklessness, and Negligence,” *Ohio State Journal of Criminal Law* 9(2): 545–582. See also Gideon Yaffe, *The Age of Culpability: Children and the Nature of Criminal Responsibility*, forthcoming, Oxford University Press.

⁴ See *United States v. Willis*, 277 F.3d 1026, 1032 (8th Cir. 2002); *United States v. Delreal-Ordones*, 213 F.3d 1263, 1268 (10th Cir. 2000); *United States v. Puche*, 350 F.3d 1137, 1149 (11th Cir. 2003).

⁵ Such a requirement is explicitly adopted by the Supreme Court in the most recent willful ignorance case to come before it. See *Global-Tech Appliances, Inc. v. SEB S.A.*, 131 S. Ct. 2060.

⁶ See Model Penal Code §2.02(7).

⁷ The Willful Ignorance Principle has always been important—plenty of people explicitly omit looking into the parcels that they are paid to carry across borders—but its prominence may grow in light of a recent Supreme Court case. In *Rosemond v. United States* (134 S.Ct. 1240), the Court ruled that knowledge about one’s confederates often suffices for complicity with the confederates’ crimes in circumstances in which many courts previously demanded some stronger commitment on the part of the defendant, such as that involved in intent. Under the Court’s ruling, Rosemond did not need to intend that one of the other people involved in their joint drug deal have or use a gun to be complicit in the gun crime that person perpetrated; it was enough that he knew that the confederate had the gun, or knew that he would use it. In light of this, the Willful Ignorance Principle provides a powerful prosecutorial tool: it allows the conviction of accomplices who were ignorant of various facts about those they aid provided that they failed to inquire about those facts earlier. If Rosemond failed to know that his confederate had a gun, he could still be complicit in the gun crime provided that he was ignorant due to a failure to inquire earlier.

Whenever we make knowledge necessary for criminal liability, the Willful Ignorance Principle licenses criminal liability also for the willfully ignorant. And so expansion of accomplice liability to a class of knowing agents is an expansion also to a class of ignorant agents, namely those who are ignorant thanks to the fact that they failed to inquire. Given the ubiquity of motives for not inquiring about those with whom we interact and aid—there are powerful social norms supporting respect for the privacy of our peers and associates, not to mention deferential presumptions regarding the good wills of our mothers and aunts—we can expect the *Rosemond* decision, together with the Willful Ignorance Principle, to result in substantial expansion of complicity liability. Imagine, for instance, that Heredia had not been driving the car, but had been a passenger instead. The *Rosemond* decision, together with the Willful Ignorance Principle, would imply that Heredia, in that hypothetical case, has the mental state of an accomplice to the importing of marijuana. In fact, if a prosecutor decided to pursue it, Heredia’s *daughters* might be charged with such a crime (most likely in juvenile court), provided that they too developed suspicions that they did not pursue prior to taking the ride across the border.

crimes and then go on to show that those same features are present in the willfully ignorant. The conditions under which the willfully ignorant can be inferred to possess the same relevant features as the knowing are the conditions under which the Willful Ignorance Principle is justifiably applied. So, although the narrow goal here is to identify when the Willful Ignorance Principle is and is not appropriately used, the broader goal is to forward a theory of the bearing of cognitive mental state on culpability that can be used to investigate the justifiability of any *mens rea* substitution principle.⁸

To tackle these questions, it will be useful to have an abstract description in hand of the agent in cases in which the Willful Ignorance Principle would seem to apply. Let's assume that at time t2 our defendant, D, performed an act A that promised him, with certainty, a benefit, B. At an earlier time, t1, he judged the probability of condition X being in place at t2 to be p. This is significant because if X was in place when he acted at t2, then A promised a harm, H, to others.⁹ And, in fact, X was in place at t2 and so H materialized, thanks to D's A-ing. At t2, the time that he A'd, D did not know that X was in place. (Perhaps at t2 D was aware that the probability of X was p; but perhaps not.) At t1, when D was aware that the probability of X was p, he could have inquired further and reached the conclusion either that the probability of X was 1, or that it was 0; he could have settled the matter through inquiry. But he did not inquire at t1 nor at any subsequent time prior to action at t2. And, finally, let's assume that D is charged with the crime of A-ing while knowing X.

Cases meeting this description differ from one another in the explanation for why D did not inquire at t1. Some agents fail to inquire to avoid prosecution for the knowledge crime. Some agents fail to inquire because they are lazy or because their minds are on other things. Some agents fail to inquire because inquiry would be costly or risky, or because somebody convinces them that this is so. Some agents fail to inquire because they are asleep at t1 or because they are knocked unconscious by someone who wants to prevent inquiry. And so on. The aim is to lump all cases involving the absence of inquiry, regardless of the explanation for the failure to inquire, and then to see how that explanation matters to the question of whether to treat D as if he knew X at t2 when he acted. Or, put another way, a characterization of the conditions under which the Willful Ignorance Principle is justifiably employed consists of a description of the class of explanations for the failure to inquire in light of which D shares the relevant features with the agent who acted knowingly at t2.

Included in this abstract description, then, are all of the following kinds of cases: (1) D is charged with knowingly transporting cocaine across the border when he did not know that he had cocaine and earlier failed to open the cocaine-filled suitcase in his trunk. (2) D is charged with knowingly soliciting sex from a minor when he did not know that the person he solicited was a minor, but also failed to inquire about her age in advance of soliciting her. (3) D is charged with knowingly receiving stolen property when he receives property that he did not know was stolen but also earlier failed to ask where the seller got

⁸ I limit myself here to the bearing of *cognitive* mental state on culpability. I believe that *volitional* mental states, such as intention, bear on culpability for related but different reasons. For some albeit incomplete discussion, see Gideon Yaffe (2004), "Conditional Intent and *Mens Rea*," *Legal Theory* 10(4): 273–310; Gideon Yaffe (2010), *Attempts: In the Philosophy of Action and the Criminal Law*, Oxford: Oxford University Press, especially Chapters 1 and 2; Gideon Yaffe (2014), "Criminal Attempts," *The Yale Law Journal* 124(1): 92–156.

⁹ No weight is being placed here on the idea of "harm." It is simply a less cumbersome term than "violation of a legally protected interest." For our purposes, for instance, importation of 350 lb of marijuana will count as a "harm," even if nobody is actually harmed by it.

it. (4) D is charged with knowingly breaking V's leg when the brakes on his moped failed, even though he did not know the brakes were faulty but also failed to check them earlier. (5) D is charged with knowingly aiding her husband in the sexual molestation of their daughter when she failed in her duty to prevent him from doing so even though she had no idea what he was doing to the daughter, but also earlier failed to knock on the daughter's locked bedroom door. What we want to know are the conditions under which it is justifiable to appeal to the Willful Ignorance Principle in support of conviction in cases such as these.

In Sect. 2, I consider what I call the “core cases”: I describe a class of willfully ignorant actors and describes the structure of the reasons that explain their failure to inquire at t_1 . In Sect. 3, I introduce a theory of the relevance of mental state to the justification of the imposition of criminal liability and punishment. In Sect. 4, I explain why, in light of the theory offered in Sect. 3, core cases of willful ignorance of the sort described in Sect. 2 are properly treated as though they involved knowledge. Sect. 3 offers, that is, a justification of the Willful Ignorance Principle, when applied to core cases. Sect. 5 reflects on the rationale for the Willful Ignorance Principle offered in Sect. 4 and in it I argue that, despite the fact that the rationale offers only a sufficient, and not a necessary condition for the justifiable use of the principle, the class of cases covered is quite a bit larger than one might have thought. The section also specifies conditions under which the Willful Ignorance Principle is properly applied in a large set of non-core cases, including both those in which the method of inquiry produces false positives and false negatives, and those in which inquiry is not costless to the defendant but could even be, itself, quite dangerous.

2 The Core Case: Omitting Inquiry to Prevent the Efficacy of One's Conscience

Consider the following description of the psychology of the defendant at t_1 , the time at which he is to decide whether to inquire about the presence or absence of X. Let's imagine that D knows himself well enough to know that if he discovers that X, then he will not A at t_2 . His conscience will interfere and lead him to refrain; his conscience won't let him impose the harm on others that he would know the act to promise. So, he anticipates that if he knows at t_2 that X, then he will not receive the benefit, B, of A-ing. Given this, he recognizes at t_1 that to inquire, and thereby to resolve the question of whether X holds, is to condition his receipt of the benefit of A-ing on the answer: he will receive that benefit, B, only if the inquiry shows that not-X, for it is only in that circumstance that his conscience will permit him to act. By contrast, he recognizes at t_1 also that if he *does not know* at t_2 that X, but instead believes at that time that the probability of X is p (or lower), as he believes at t_1 , then he will A at t_2 , and so reap the benefit, B, of A-ing, although he will also impose a risk of magnitude p of harming others. (The risk of the harm in acting at t_2 is p since that is the risk that X will be in place at t_2 ; or so he believes at t_1 .) If he inquires, his chance of receiving B equals the chance that X does not hold; if he does not inquire, then he is sure to receive B, but the chance that he will also inflict H is equal to the chance that X holds. So, in deciding whether to inquire, he is deciding whether to allow his receipt of the benefit to be conditional on the presence or absence of the condition that guarantees harm. Or, put another way, in deciding whether to inquire, he is deciding whether to face a risk of magnitude p that he will not receive the benefit, on the one hand, or impose a risk of magnitude p that others will suffer the harm, on the other.

There are two further identifying features of core cases: first, inquiry is intrinsically costless. At t_1 , D would not have had to pay a firm to do a study to determine whether X was in place, nor would inquiring itself have risked harm to himself or to others that was independent of the outcome of the inquiry, nor would inquiry itself have subjected him to police investigation that would have been either dangerous or costly for him.

The second feature of core cases worth highlighting is this: in core cases, D believes at t_1 that inquiry will be successful. That is, D believes that if X will, in fact, be present at t_2 , then inquiry at t_1 will lead him to believe that X will be present at t_2 ; and, if X will be absent at t_2 , then inquiry at t_1 will lead him to believe that X will be absent. Inquiry's false positive and false negative rates are both zero. This need not be *true* for the case to be a core case, but it must be believed by D at the time at which he omits inquiry. In Sect. 5, I will show that, in a large class of cases (although not all) in which one or both features are absent—inquiry is believed to be costly and sometimes yields false positives or false negatives—the Willful Ignorance Principle can still be safely applied for the same reasons that it is in core cases.

Between this section and the next two, I will argue that, if D fails to inquire at t_1 because, when looked at as just described, the cost-benefit analysis comes out in favor of omitting inquiry at t_1 , then the Willful Ignorance Principle is justifiably applied in his case; he deserves to be convicted of knowingly A-ing at t_2 , even though he did not know X at t_2 .

We can summarize how things look, in a core case, to the agent at t_1 with the Table 1.

An explanation of Table 1 is in order. There are four basic conditions represented here, depending on whether D acts or refrains and depending on whether X is in place or not. It is a 2×2 matrix in which each of the four boxes is divided into three sub-boxes. The outcome rows indicate what benefits and harms accrue in each of the four conditions: D receives B if and only if he acts at t_2 ; and the harm, H, is inflicted if and only if D acts and X is in place. The conditional probabilities are the probabilities of reaching the relevant condition of the four, given inquiry at t_1 , or given the absence of inquiry at t_1 . Put another way, to calculate the relevant conditional probabilities, we assume that D inquired at t_1 (or failed to inquire) and we ask ourselves, given that presumption, what the probability is that he A'd with X present, A'd with X absent, refrained with X present, and refrained with X absent.

Table 1 Represented are (1) the four outcomes if X holds or fails to hold, and if the agent acts or refrains at t_2 , and (2) the conditional probabilities of the relevant condition obtaining (e.g., X/Act or not-X/Refrain), given inquiry or given the absence of inquiry at t_1

		X	Not-X
Act at t_2	Outcome	B & H	B & not-H
	Prob Inquiry at t_1	0	1-p
	Prob No Inquiry	p	1-p
Refrain at t_2	Outcome	not-B & not-H	not-B & not-H
	Prob Inquiry at t_1	p	0
	Prob No Inquiry	0	0

The resulting conditional probabilities in the table need some more detailed explanation. First, consider the Act/X box (the upper left of the four). If D inquires, then he will discover that X and so he will refrain with probability 1. So, the probability that D ends up in that condition, given inquiry, is zero. By contrast, the probability that D ends up in that condition, given the absence of inquiry, is just the probability of X, namely p , since D is guaranteed to act if he does not inquire. Move now to the Act/Not-X box (the upper right). If D inquires, he will discover that not-X, and so he will act in order to reap the benefit of acting. So, the probability that he will end up in that condition is the probability of not-X, namely $1-p$. Similarly, the probability that D will end up in the Act/Not-X box if he does not inquire is $1-p$. Since D is guaranteed to act in the absence of inquiry, that probability just equals the probability that not-X. When we move to the Refrain boxes (the lower left and right), we see that the probability of D ending up in either of those boxes in the absence of inquiry is zero, since the probability of refraining in the absence of inquiry is zero regardless of whether X or not-X; he will act unless his conscience prevents him, and that will not happen if he does not inquire. By contrast, the probability of ending up in the Refrain/X box (the lower left) is equal to the probability of X ($=p$), given inquiry, since if D inquires and finds that X, then he will refrain with probability 1. And, correlatively, the probability that he will end up in the Refrain/Not-X condition (the lower right) is zero, given inquiry, since if he discovers that not-X, then he will act and so will refrain with probability zero. As a check on this analysis, note that the four values, given inquiry, should add up to 1, as should the four values given the absence of inquiry; and they do.

In the core case, then, the agent's failure to inquire at t_1 about X derives from his recognition that an inquiry might show him to have decisive (altruistic) reason not to act at t_2 . This means that inquiring risks closing off the possibility of receipt of the benefit of acting. Further and importantly, that risk is of precisely the magnitude that he takes the risk that X will hold at t_2 to be. After all, to believe that there is a probability of p that X necessarily involves the belief that the probability is p that an inquiry will establish X, provided (as in core cases) that inquiry will yield neither false positives nor false negatives.

All of the examples briefly described above could be examples of the core case. The agent might not look in the suitcase in his trunk, before he drives across the border because he knows that, if he finds cocaine there, he will get cold feet and thus fail to receive the \$1000 that he has been offered to deliver the suitcase. He might not ask about the age of the child he is about to solicit for sex because he believes that, if he knows she is a child, thoughts of his own children will preclude him from acting. And so on. This is not to imply that all such cases are examples of the core case, but only that many that meet the sparse descriptions of such cases, given in the Introduction, when more details are uncovered, turn out to be instances of the core case. The question is just whether the failure to inquire is explicable through appeal to the information about the various outcomes contained in Table 1. If it is, then we are looking at an instance of the core case.

3 The Point of *Mens Rea*

How does the reconstruction offered in Sect. 2 of D's point of view at t_1 , when his is a core case, bear on his culpability for A-ing at t_2 ? To answer that, we need a tool that bridges the gap between the probability-corrected values associated with acting and refraining that the agent anticipates at a time of decision and his level of culpability for the decision. The bridge to be offered in this section derives from the idea that we care about an agent's mental states, and

the deliberative processes they guide, when assessing his responsibility because, thanks to them, his actions manifest his culpability-relevant values. In particular, thanks to the agent's mental states, his actions manifest *the evaluative weight that he gives to his own interests in comparison to the interests of other people*. When a tendency to put himself first, to a greater degree than is acceptable, is manifested in his conduct, he is criminally culpable for that conduct.¹⁰ For ease of exposition, call this the agent's "social preference."¹¹

When determining what social preference is manifested in an agent's behavior, such as the behavior of bringing drugs across the border, we are constrained by a presumption of rationality and a principle of lenity. The presumption of rationality requires that we take the agent's behavior to be genuinely reflective of his values. We assume, that is, that he is not acting in a way that is, by the agent's own lights, insufficiently supported by reasons. The principle of lenity demands that we interpret his behavior to manifest a social preference as close as it can be reasonably taken to be to an acceptable one. It demands, that is, that we depart as little as is reasonably possible from the conception of him as someone who cares as much about others as he ought. We find criminal culpability when, even under these two constraints, the agent's conduct manifests an unacceptable social preference.

This line of thought leads naturally to a conception of criminal culpability assignments bearing a close similarity to Hand Formula assessments of negligence.¹² Consider the following simple equation, where EV_{Self} is the expected value for the agent himself and EV_{Other} is the expected value for others¹³:

$$\text{The value equation : Value (action)} = EV_{\text{Self}} (\text{action}) + a * s * EV_{\text{Other}} (\text{action})$$

The a parameter is a measure of the agent's attention to the potential impact of his act on other people. If he does not attend to that at all—if $a = 0$ —then his valuation of the act is uninfluenced by his judgments concerning the impact that that act will have on others. When $a = 0$, that is, it is not that he fails to make judgments about what effect his conduct will have on others; it is, rather, that the judgments that he makes are not manifested in his behavior.

¹⁰ For more about very closely related ideas, see, as a start, John Fischer and Mark Ravizza (1998), *Responsibility and Control: A Theory of Moral Responsibility*, Cambridge: Cambridge University Press; Gideon Yaffe (2012), "Intoxication, Recklessness, and Negligence," *Ohio State Journal of Criminal Law* 9(2): 545–582; Gideon Yaffe (2010), *Attempts: In the Philosophy of Action and the Criminal Law*, Oxford: Oxford University Press; Peter Westen (2006), "An Attitudinal Theory of Excuse," *Law and Philosophy* 25(3): 289–375; Thomas Scanlon (2000), *What We Owe to Each Other*, Cambridge: Harvard University Press; Pamela Hieronymi (2006), "Controlling Attitudes," *Pacific Philosophical Quarterly* 87(1): 45–74; Angela Smith (2005), "Responsibility for Attitudes: Activity and Passivity in Mental Life," *Ethics* 115(2): 236–271.

¹¹ The idea of social preference is intended to capture at least three distinct facts about an agent. An agent who cares less about others than he ought to might take the fact that his act will cause harm to others to give him no reason whatsoever to refrain from it. That would be a failure to properly *recognize* legal reasons for action. Or he might recognize that that fact provides a reason to refrain, but grant less reason-giving weight to that fact that he ought to. This would be a failure of *weighing*. Even if he does grant proper weight to the fact that an act he is considering causes harm to others, he might not respond to it in his deliberations as he ought. This is a failure of *response*. For our purposes here, this tri-partite distinction, and the subsequent variation in possible cases one finds, need not detain us.

¹² According to the Hand Formula, a tort defendant's failure to take a precaution against plaintiff injury is "unreasonable" provided that the burden (B) of taking the precaution is less than the expected loss to the plaintiff thanks to such failure (PL, or the product of the probability of the loss and its magnitude).

¹³ I developed this idea in discussion with Antonio Rangel. I have presented it and considered its implications for *mens rea* also, in Gideon Yaffe (2012), "Intoxication, Recklessness, and Negligence," *Ohio State Journal of Criminal Law* 9(2): 545–582.

If the agent cares not at all about others—if $s = 0$ —then he will grant no weight to the impact of his act on others in his valuation of the act, even if he is vividly aware of what that impact will be (even if, that is, a is high). And, conversely, if he is perfectly egalitarian in his weighing of others' interests—if $s = 1$, implying that he weighs others interests exactly as heavily as his own—then he will grant the impact of his act on their interests the very same weight as he grants to its impact on his own, in his valuation of the act, provided he is aware of what the impact of his act will be on others (provided, that is, that a is high).

The presumption of rationality implies that the agent's act manifests this fact about the agent: Value (action) > Value (not-action). That is, the presumption of rationality mandates that the agent acted as he did because he valued such action over refraining from it. When we use this fact, together with information about the expected values of the act for the agent and for others, we are able to assign an upper-bound to the product $a * s$.

Information about the agent's psychological state at the time of action gives us information about the value of the a parameter. The principle of lenity mandates that we assign a value to a that is as low as possible, consistent with the evidence. After all, the lower a is, the higher s can be consistent with the fact that Value (action) > Value (not-action). For now, we will assume that we know enough about the agent at the time of action to conclude that he was vividly aware of the potential harms to others his action would cause; or, rather, that he was no less aware of them than he was of the benefits to himself that his action might cause. As we will see, when information about an agent places this assumption in doubt, then it can also make it inappropriate to employ the Willful Ignorance Principle. But, for now, we will assume that $a = 1$.

Ultimately, it is the value of s that we are interested in when we are assessing culpability. We want to know how heavily the agent weighed others' interests both intrinsically and in comparison to his own when he acted in a way that was, in fact, in violation of their legally protected interests. If, given the presumptions of lenity and rationality, we are able to infer that he granted next to no weight to others' interests in comparison to his own—if, for instance, the potential harm to others from the act was great, the potential benefit to himself was negligible, and he was vividly aware of the potential harm—then we are inferring that the agent's s value was very low. That is, his was a highly culpable act.

Consider, for instance, someone who burns down a building, knowing it to be occupied, for the insurance money, and thereby kills the occupants. To say that he is in a state of knowledge with respect to the harm to the occupants (setting aside the harm to the insurance company and its customers) is to imply two things: (1) that the agent's a level was high—he was aware of the potential harm to those in the building when he acted—and (2) he took it to be virtually certain that those inside would burn when he acted and so EV_{Others} was catastrophically low. By contrast, to say that he was consciously aware of a risk is to imply that his a level was high, but it is also to imply, in contrast to the case of knowledge, that EV_{Others} was not as low, since he thought there was some chance that the harm to others would not materialize. This is the crucial psychological difference between the knowing and the reckless actor. And this crucial psychological difference permits a different inference about culpability: given the difference between the knowing and reckless actors' probability assignments, and given that they both have high a values, and given that the results for the two actors and for others are the same, and given presumptions of lenity and rationality, a higher s value is manifested in the actions of the reckless than of the knowing actor. He cares more about others' interests than the knowing actor and so is less culpable for his bad behavior (although culpable enough to warrant some criminal penalty).

The point just made can be expressed graphically as Fig. 1.




Fig. 1 $\text{Prob}(\text{result_for_others}|\text{act})$, on which EV_{others} depends, is on the x-axis. The s value is on the y-axis. The highest s value that can be inferred for each value of $\text{Prob}(\text{result_for_others}|\text{act})$ (given that $a = 1$, the other parameters of the Value Equation are set, and the Value Equation > 0 as implied by the presumption of rationality), is represented in the curved line. A legislative decision to set the *mens rea* standard at recklessness is the decision to set a criminal penalty, given an s value at or below the higher of the two dotted lines. To set the *mens rea* standard at knowledge is to choose the s value that can be inferred when $\text{Prob}(\text{result_for_others}|\text{act}) = 1$ as required for culpability

Two crucial points are illustrated by Fig. 1. First, *ceteris paribus* the reckless agent—who is aware of a certain probability of the harm to others that falls short of 1—is less culpable than the knowing agent, who is aware of a probability of approximately 1 of this potential harm. This result falls out of the theory of *mens rea*'s point offered here. The reason is that the reckless agent's s level manifested in his action (given the assumption that he takes the value of the act to be greater than zero) is below a higher threshold than the knowing agent's can be inferred to be below. Given a presumption of lenity—we hold agents responsible for being the best they can be consistent with what we know of their behavior and their mental states—knowing agents manifest in their behavior worse social preferences than do reckless agents. Second, the greater the probability that the agent takes there to be of harm, the lower is the threshold that we can infer his s value to be below. At least, this is so given the assumption that his a value is 1—that is, given the assumption that he is vividly aware of the risks when he acts.

The view that will guide our discussion from here forward, then, is that mental states matter to culpability because they allow us to make educated guesses about crucial pieces of information appealed to in the Value Equation—namely, information about the value of the a parameter, and the expected values for the agent and others of the action. These educated guesses, together with the constraints of rationality and lenity guiding our usage of the Value Equation, allow us to reach a conclusion about the social preference manifested in the agent's act and summarized in the s value. Since to manifest a problematic

social preference is to be criminally culpable, we have in hand an account of the point of *mens rea*.

4 Justifying the Willful Ignorance Principle

With the picture of the point of *mens rea* offered in Sect. 3 in mind, let's return to D, who did not inquire at t1 about X; performed A at t2; thereby provided himself with benefit, B, and others with harm, H; did not know at t2 that X was in place; and is charged with the crime of A-ing while knowing that X. And let's assume that D's is a core case: inquiry was costless and perfect (the method yields neither false positives nor false negatives), and he failed to inquire at t1 because inquiry would have risked the possibility that he would discover X, a discovery that would prompt him to refrain from acting at t2, and would thus preclude the possibility of his receiving the benefit of such action. Using the framework just described, and referring to Table 1, consider what *s* value was manifested in D's failure at t1 to inquire about X. When D failed to inquire, how seriously was he weighing harm, H? Did he weigh it less heavily than he weighed the impact of his conduct on himself? The same?

Recall that, as we are now reconstructing D's psychology at t1, he sees inquiry as assuring that there will be no harm to others while risking a failure to receive the benefit of action; and he sees the absence of inquiry as assuring receipt of the benefit while risking harm to others. Given that D did not inquire about X, the presumption of rationality implies that he took there to be greater reason not to inquire than there was to inquire:

$$\text{Value (not-inquire)} > \text{Value (inquire)}$$

This implies, under the Value Equation, and given the various assumptions we have been making:

$$(B * 1 - 1 * s * H * p) > (B * (1 - p) - 1 * s * H * 0)$$

Put in English: we are able to infer from the decision not to inquire that the agent takes the scenario in which he is sure to get the benefit, but might or might not cause the harm, to be of greater value than the scenario in which he might or might not get the benefit, but will avoid the harm.¹⁴ Simplifying algebraically, we reach the following:

$$B - s * H * p > B(1 - p)$$

Solving for *s*, we reach the following result:

$$s < B/H$$

If B is of smaller magnitude than H—if the potential benefit to oneself of the act is smaller than the potential harm to others—then we learn from D's decision not to inquire that he cares less about others than he cares about himself. If, for instance, the act promised him a potential benefit of 5 dollars, and promised a potential loss to others of \$10, we can learn from his decision not to inquire that his *s* value is below one half. He cares twice as much about the impact of the act on himself as about its impact on others.

¹⁴ Attentive readers will notice that a “+” in the Value Equation has become a “−” here. This is because H is a positive value—the magnitude of a harm—and it is assumed, consistent with the lenity constraint, that *ceteris paribus* D takes the prospect of H's occurrence to count against performance of the action rather than for it.

Now let's imagine that, at time t_2 , when the agent acts after having earlier failed to inquire about whether X would obtain at t_2 , he is consciously aware that the probability that X obtains is p . So, as we are assuming about his psychological state at t_1 , at t_2 $a = 1$ and $EV_{\text{Others}}(\text{action}) = H * p$. From the fact that he acted, we can conclude the following:

$$\text{Value}(\text{act}) > \text{Value}(\text{refraining})$$

Using the Value Equation, we reach the following result:

$$B * 1 - 1 * s * H * p > 0 * 1 - 1 * s * 0 * 1$$

Solving for s , we reach the following:

$$s < B / (H * p)$$

Note that, since p is between 0 and 1, under the most lenient assumption, the s value that can be inferred solely from D 's action at t_2 , while aware of a risk of p that X obtained, was higher than the s value that we could infer from his decision not to inquire. That is, if we did not know that he failed to inquire, we would conclude that his s value is greater than the ratio of the magnitude of the benefit and the harm, while from his decision not to inquire, we are able to infer that his s value is at most that ratio. This result alone helps to show how the decision not to inquire bears on the agent's culpability. Assuming that there has not been a change in the agent's s value between t_1 and t_2 , we know that the agent was more culpable for his act given his decision not to inquire than he would have been in the absence of that decision.

But, in fact, we can say more even than just this. Imagine that at t_2 the agent *knew* that X obtained and went ahead with the act anyway. Imagine, that is, that all is as just described, except that at t_2 he takes the probability of the harm to be 1. What could we conclude about his s value? Well, substituting 1 for p in the result just reached we find the following:

$$s < B / H$$

That is striking, for that is precisely the ceiling on s that could be inferred from the decision not to inquire. So, what have we learned? We have learned that, under the various assumptions that we have made so far, we can infer that the s value of the agent who does not inquire is below a certain threshold; and we can infer that the s value of the agent who acted with knowledge is below a certain threshold; *and the two thresholds are the same*. Add the view described above, according to which mental state matters to culpability because of what it evidences about the degree to which the agent cares about his interests in comparison to others, and we learn the following: *the willfully ignorant agent is precisely as culpable as the knowing agent*. This provides a rationale for the Willful Ignorance Principle. Provided that the defendant fails to inquire so as to prevent his conscience from later keeping him from acting, the facts that matter to the justification of treating knowing agents as we do are present, also, in his case. Knowledge matters because of what it says about the agent's s value; and the very same thing is said about the agent's s value by willful ignorance, provided that the agent's psychology is an example of the core case.

The argument just offered can be seen clearly by reflecting on the Fig. 2.

Although it is hard to specify with any precision, there is some level of s that we demand of citizens. There is some degree to which people must care about particular forms of impact of their conduct on others when they act. Below this level we have a potentially criminal mind, and so potentially criminally punishable conduct. What we see in this

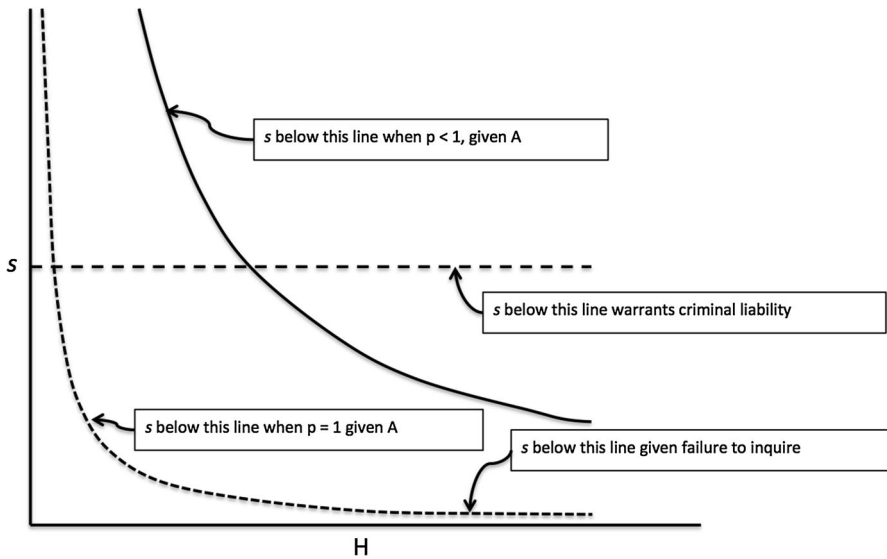


Fig. 2 Values of H increase along the x -axis. Assuming a given level of benefit, B , from the act and given levels of probability of benefit and harm, as H increases, s , which increases on the y -axis, can be inferred to be lower. What is depicted here are three thresholds of s that can be inferred, given various single pieces of information: (1) $D A'd$ at t_2 while believing p to be less than 1, (2) $D A'd$ at t_2 while believing p to be 1, and (3) $D A'd$ at t_2 after having failed to inquire about X at t_1

chart is that, as the potential harm to others of the act at t_2 increases, we are able to infer a lower level of s on the part of the agent. We also see that we are able to infer lower levels of s when the agent knows that X —when his belief that the probability of X is 1—in comparison to the level that we are able to infer his s level to be when he takes there, instead, to be a risk of X —when his belief that the probability of X is less than 1. Further, and of central importance here, is that, when he fails to inquire at t_1 for the reasons outlined here as constituting the core case, we are able to infer that his s level is below precisely the same threshold as that that can be inferred from knowing action. It is this coincidence in thresholds that justifies applying the Willful Ignorance Principle in core cases.

5 The Bounds of the Willful Ignorance Principle

To identify a class of cases in which the Willful Ignorance Principle is justifiably applied, as we have in the last few sections, is to identify a sufficient condition, but not a necessary condition, under which it is appropriately applied. But this raises the question of whether there are non-core cases in which we can apply the Willful Ignorance Principle. Under what conditions is a non-core case nonetheless similar in relevant respects to a core case? My primary aim in this section is to identify such conditions. Towards that end, however, several observations about the application of the Willful Ignorance Principle in core cases are warranted. By understanding the core case better, we will be positioned to understand what features of a core case must be preserved if the Willful Ignorance Principle is to be safely applied.

5.1 Some Observations about Willful Ignorance in the Core Case

5.1.1 *Self-Flattery and Nerve*

Notice that there is something somewhat paradoxical about the psychology of the agent in the core case. It is because the agent believes himself to be such that if he knows that X he will not act that he does not inquire about whether or not X is so. And we have shown that, when he acts from that belief, he shows himself to have as low an *s* value as someone who knows that X and acts anyway. It follows that, if his *s* value at *t*₂ is the same as his *s* value at *t*₁, when he omits to inquire, then he is motivated by a belief *that is shown to be false by the actions he performs in light of that belief*. Only someone callous enough to act in the face of knowledge that X would refuse to inquire about whether X holds in order to prevent himself from soft-heartedly refraining later and thus denying himself the benefits of acting. In short, he would be motivated not to inquire only if he has no reason to fear that inquiry would place his benefits at risk. If his conscience would actually hold him back from acting, it would also hold him back from omitting inquiry.

While this is all true, it is important to note that there is nothing problematic about it for the view proposed here. We cannot conclude from the fact that the agent acted at *t*₂ while aware of a risk that X that his *s* value is as low as a knowing agent's. But we can reach that conclusion, given the prior failure to inquire. That the inference depends upon false consciousness on the part of the agent—he thinks better of himself than, it turns out, he ought to—in no sense weakens the inference. In fact, it is arguably a benefit of the view proposed here that it paints the agent in the core case as suffering from a deluded sense of his own virtue. There is an intuitive sense, I believe, in which willfully ignorant agents, deserving of treatment as if they were knowing, think better of themselves than they ought. There is something false or phony about their pleas to ignorance. This is, on its face, peculiar since *they are ignorant* and so, what they say when they note that they did not know the relevant facts is perfectly true. What is phony about their pleas is that they are saying that their ignorance paints them in a better light, when, in fact, it does not. This is so despite the fact that it is, undeniably, true that they are ignorant when they act.

What this shows, however, is that included in the class of core cases are not just those cases in which the agent falsely believes that his conscience will prevent him from acting, but also those cases in which the defendant's belief is true. These will include all cases in which the defendant thinks that knowledge that X is in place will impede action because he knows that he won't have the nerve to act in the face of knowledge. In fact, in some such cases, what the agent experiences as his conscience interfering is actually a failure of nerve. Some agents, for instance, will probably feel, when they know that X and so cannot act, that they are refraining from action due to a recognition of A-ing as wrongful, given X. But, in fact, they are simply too frightened to A, given X, and cover their fear with a self-flattering story of the restraint of conscience. Either way, however, what matters to the applicability of the Willful Ignorance Principle in the core case is the belief at *t*₁ that knowledge will impede action at *t*₂, whether that belief is true, and whether, when true, it is true for the reasons that the agent would like to believe it to be.

5.1.2 *D's t₁ Psychology*

Nothing about the explanation for the propriety of applying the Willful Ignorance Principle in core cases depends directly on any claims about the psychology of the agent at *t*₂. So,

the class of core cases includes agents who differ widely in their psychological states at that later time. While it is important that *at t1* the agent's *a* value is high—he must be aware of the potential harms of the *t2* act for it to be true that he values omitting inquiry more than inquiring—he could be completely oblivious to the potential harms of the act *at t2*, when he actually acted, and still be rightly convicted of a knowledge crime. What this implies is that the class of core cases includes those in which the agent is in a state of mind at *t2* that makes him culpable for some crime less bad than a knowledge crime, but also includes those cases in which the defendant is in an entirely blameless state of mind at *t2*.

Although D's *t1 a* value must be high, there are few other constraints on D's *t1* psychology that must be met for the Willful Ignorance Principle to be applicable in a core case. Thus, the characterization here of the core cases of willful ignorance shows them to be inaptly labeled as “willful.” There are many instances of the core case in which D does not *will* to omit inquiry at all, much less that he omits inquiry intentionally. What is crucial is that he omits inquiry, thanks to the fact that inquiring would place the benefit of acting at risk. His failure to inquire must be *explicable* through appeal to the anticipated benefits, harms, and risks summarized in Table 1.

But to say that the omission is *explicable* in this way is not to imply that the agent engages in a psychological process in which he calculates the value of inquiring and not inquiring, calculations which lead to a decision to omit inquiry. It is possible that, in light of his beliefs about those benefits, harms, and risks, he is able to do what makes most sense to him—namely omit inquiry—without ever expending energy on such a deliberative psychological process at all. It might be that the answer to the question “Why did he omit inquiry?” is “Value (not inquiry) > Value (inquiry)” without it also being true that he *had the thought*, conscious or unconscious, that this was so.

Because this last point is essential to the case for the claim that the class of core cases is very broad, it is worth making the point again in a slightly different way. Imagine that you were designing a system that would omit inquiry whenever Value (not inquiry) > Value (inquiry). It could be that there is a large class of cases in which this is either so obviously true, or so likely to be true, that there would be no need for the system to make the explicit calculation. Perhaps there are signs of the truth of this inequality that are sufficiently reliable that it is more efficient for the system to just omit inquiry when those signs are present than to expend energy calculating the values of the alternatives, much less comparing them. The designer of the system would, then, do best to allow the presence of the sign to prevent inquiry without performing the calculation. In that case, the failure to inquire would be *explained* by the fact that the value calculation came out a certain way—the system would not respond in the way it does if this were not the case—without it being also true that the system calculated the quantities appealed to in the explanation.¹⁵ In the same way, it would be possible for a human agent's failure to inquire to be explicable by appeal to his beliefs about the benefits, harms, and probabilities of the outcomes described here without engaging in any psychological process in which the relevant value calculations are made.

Similarly, one can, in theory, imagine devising a system that acts in a certain way when possessing a series of pieces of information *because they add up a certain way* even if the system never adds them up. A trumpet makes a certain tone when the first and third valves are closed. To cause the trumpet to make that sound, the player simply closes both valves while blowing on the trumpet. The player need not do yet something further to combine the

¹⁵ The line of thought here intersects with the view of omission I presented in Gideon Yaffe (2016), “In Defense of Criminal Possession,” *Criminal Law and Philosophy* 10(3): 441–471.

two acts of pressing the valves in order to cause the trumpet to produce the tone. And, similarly, it would be possible for a creature to be designed such that when the variables in the Value Equation are assigned, it acts in a way that is predicted by the Value Equation even if the creature never actually produces a new representation with the content that would result from application of the Value Equation.

The existence of the explanation for the failure to inquire in core cases is what supports applying the Willful Ignorance Principle in such cases. While that explanation will be available in many cases, thanks to the explicit thought process that leads D to omit inquiry at t_1 , there will be other core cases in which the explanation applies for subtler reasons. What this implies is that we should not think that what we are doing when we assess willfully ignorant agents is to imagine that the failure to inquire is itself a crime for which we assess culpability; we are not holding the willfully ignorant agent responsible for failing to inquire. If we were, then in cases in which the failure to inquire was not willful, in the ordinary sense, there could be no culpability. But, since the failure to inquire need not be willful for the Willful Ignorance Principle to apply, it cannot be that we are holding the agent responsible for his failure to inquire. The only crime is knowingly A-ing. Failing to inquire at t_1 is entirely insufficient for the crime. What is sufficient is A-ing at t_2 in a complex set of circumstances, including a failure to inquire at t_1 , thanks to which the agent's s value at t_2 , manifested in his act, is the same as that of the knowing actor.

5.1.3 *A's Manifestation at t_2 of D's t_1 Values*

The Willful Ignorance Principle is not safely applied in a core case unless the t_2 action expresses the very social preference that could be inferred from D's t_1 failure to inquire. If that morally relevant feature of the agent is not manifested in the conduct at t_2 , then that conduct is not criminally culpable. It is unjustifiable to punish bad people for bad behavior unless their bad behavior manifests what is bad about them. So, under what conditions does the t_2 behavior manifest the problematic social preference that is manifested in the t_1 failure to inquire?

It must at least be the case that the agent's trait of interest—namely, his social preference—is stable from t_1 to t_2 . Without this stability, no inference about the agent's s value at t_1 can illuminate his s value at t_2 , when he acts. It is easy to imagine cases in which this stability is absent. Perhaps the time gap between t_1 and t_2 is very long and D undergoes important experiences during that period, experiences that we can expect to shape his s values. Perhaps people's s values are not, in general, stable things but are, instead, context-sensitive. One might think, for instance, that adolescents have high s values when in the presence of positive influences, such as peers with high s values, and much lower s values when in the presence of negative influences, such as callous peers. If t_1 and t_2 differ with respect to these contextual factors, then, again, we are not safe applying the Willful Ignorance Principle, even in a core case. But the broader point should not be lost: to the extent that there is no reason to think that D's s values are unstable from t_1 to t_2 , it is potentially true that his t_1 social preferences are manifested in his t_2 behavior and so, potentially, it is safe to apply the Willful Ignorance Principle. While it is an empirical question how frequently we find the relevant kind of stability, a bit of experience with human nature would suggest that the domain of safety here is quite large.

But stability in morally salient features from t_1 to t_2 does not suffice to support the claim of t_2 culpability. We need more. The crucial point is that, in holding D criminally responsible for his t_2 conduct, we are judging his culpability with reference to his t_2 social preferences manifested in his action. But the t_1 failure to inquire provides us with a piece

of evidence about the manifested social preference. We are licensed to appeal to more than D's t_2 mental states when assessing his culpability, and so we are licensed to find him more culpable for his t_2 conduct than we would have found him had we examined only his t_2 mental states and behavior. While the issue here is difficult, it is not harder in the context of willful ignorance than in any other context. We frequently need to appeal to past evidence in order to make inferences about present facts that are unavailable to direct perception. And this is just as true about the assessment of social preference as it is of any other fact. Consider, for instance, the bearing of a person's threat on day 1 to kill another on day 2. That threat is surely probative of the question of whether he acted with intent to kill on day 2, or with some "lesser" mental state. This is so despite the fact that he might have changed his mind, or might have fully repented between days 1 and 2, or even if mental states are sensitive to changes of context of the kind we find between day 1 and day 2. The question of when past evidence is and is not of use is not more difficult in the context of willful ignorance than it is elsewhere.

5.1.4 Low Probability Catastrophe

The Value Equation is an inferential tool. It is useful for making inferences about the social preferences manifested in D's actions. In general, inferential tools are not to be used in circumstances in which we have significant doubt about their accuracy, especially when faulty inferences would have very bad consequences. Such is the case with the Value Equation when the agent takes the probability that his act will harm others to be very low. The reason is that estimates of what probabilities a person assigns to an event are always subject to errors of a few points, at least. When the probabilities in question are low, small absolute errors are very large errors in percentage; since inferences about s values are highly sensitive to differences in percentage of change in probability estimates, they are highly sensitive to very small absolute errors when the probabilities are low. For example: if we judge D to have thought the probability was .01 when it was actually .005, then an absolute error of only one half of a percentage point doubles the probability that D assigned. When that is the case, the inferable ceiling for the s value is extremely sensitive to small variations in the absolute magnitude of the probability of harm to others. In the example just given, the second term in the Value Equation ought to be half the size that is actually given to it when we mistakenly judge D to have thought the probability of harm was .01 rather than .005. So, this very small error can result in a very large mistake about D's s value. It could be the difference between the conclusion that D's s value is one half and the conclusion that it is 1. By over-estimating the probability of harm by 100 %, we will infer an s value half of what we are warranted in inferring. The result is that the Value Equation cannot be safely used to assess culpability when the probability of harm is very low; there is just too much potential for very large errors in the assessment of s value.¹⁶

To put the point more concretely, imagine that D makes an assessment of the probability that turning the key in his car's ignition will set fire to his neighbor's home; he judges the probability to be one in a million. Before heading to the store in his car, he omits inquiring as to whether there is a trail of oil from below his car to his neighbor's home, and omits inquiring as to whether his starter is sparking unexpectedly. Further, let's imagine his is a core case: his failure to inquire is explained by the way in which inquiry would place the

¹⁶ In Gideon Yaffe (2012), "Intoxication, Recklessness, and Negligence," *Ohio State Journal of Criminal Law* 9(1): 545–582, I make this point in support of the "substantiality" prong of the Model Penal Code's definition of "recklessness".

benefits of going to the store at risk. Sadly for all involved, the turn of the key causes the neighbor's house to burn. Is D as culpable for this as a knowing arsonist? The answer is that whether he is or not, we are not in a position *to say* that he is for we cannot infer that he granted the harm to his neighbor as little weight as the knowing agent would. The obstacle to the making of this inference is that our best tool for making it, namely the Value Equation, is not available to use in cases like this one in which the probabilities the agent assigns to the event are very, very low.

5.2 Cases “Relevantly Similar” to the Core Cases

In Sect. 5.1, I identified potential obstacles to the safe use of the Willful Ignorance Principle in core cases. This provided us with some grounds for caveat. In particular, there is reason to think that even in a core case the Willful Ignorance Principle is not safely applied if something defeats the evidential import of the t_1 s value for understanding D's t_2 s value. And it is not safely applied if the probability of harming others is very low. My goal in this subsection is to identify a few different ways in which a case can be different from the core case, and yet the Willful Ignorance Principle can be safely applied for the reasons that it is safely applied in core cases. It is important to note, however, that the same caveats apply to these “relevantly similar” cases as applied to core cases. Even in a non-core, relevantly similar case, for instance, the Willful Ignorance Principle cannot be applied if the probability of harm is very low.

5.2.1 Duties Not to Inquire

So far, I haven't said anything about the bearing, if any, of duties either to inquire or to omit inquiry on the applicability of the Willful Ignorance Principle. Ethical duties of both sorts are easy to find. Parents have duties to inquire about where and how and with whom their teenage children will be spending their time. Doctors have duties not to inquire about aspects of their patients' personal lives that do not bear on their health. Whole areas of law—think of the law governing disclosures in real estate transactions—are driven by conceptions of what people do and do not have duties to inquire about, since that bears critically on the question of what information others have a duty to disclose.

We can start with an observation about cases in which D has a duty to inquire that he shirks by omitting inquiry. This fact alone does not suffice to make the Willful Ignorance Principle applicable.¹⁷ It does not follow from the fact alone that someone shirks a duty to inquire that he is just as culpable for his later ignorant action as someone who performs such an action knowingly. It does not follow from the fact alone that D ought to have asked his teenage son where he was going that he is as culpable as someone who knowingly aided his son in selling drugs. Similarly, if the case is a core case—the father's omission of inquiry is explicable by appeal to the fact that inquiry would have placed the benefits of later action at risk, for instance—then the Willful Ignorance Principle applies, and the fact that D shirked an obligation to inquire is irrelevant to the case for that claim. In short, what

¹⁷ Alex Sarch has suggested that a duty to inquire, shirked by the omission of inquiry, is essential for the safe application of the Willful Ignorance Principle. See Alex Sarch (2014), “Willful Ignorance, Culpability and the Criminal Law,” *St. John's Law Review* 88(4): 1023–1101. However, if my argument in this paper succeeds, Sarch is mistaken. A case can be core, even taking into consideration the various caveats discussed above, without the defendant having any duty to inquire. Or, to put the point another way, an agent's failure to inquire can be explicable by appeal to the way in which inquiry would put the benefit of action at risk without the agent having any duty to inquire.

matters is whether the case is a core case, not whether the agent shirked an obligation to inquire.

Requiring subtler treatment, however, are those cases in which D has an obligation *not* to inquire.¹⁸ Recall that crucial to the applicability of the Willful Ignorance Principle in core cases is the explanation of the failure to inquire by appeal to the way in which inquiry places receipt of the benefit of action at risk. This fact allows us to reach our first, negative, result: if the failure of inquiry is explained not this way, but instead by the fact that D had a duty not to engage in the relevant inquiry, then the Willful Ignorance Principle fails to apply for the reasons that it applies in core cases.

However, there are also cases in which the failure to inquire is overdetermined: D did not inquire in part because he had a duty not to inquire, but also because inquiry would have placed receipt of the benefit of action at risk. For instance, a pharmacist gets paid for selling pseudoephedrine; the money is the benefit of action. He fails to inquire about what his customer will use it for. In part, this is because he thinks it his duty to respect his customers' privacy when it comes to matters of health. But it is also because he believes that, were he to inquire and discover that his customer is intending to use the drug to make methamphetamine, he would then refuse to sell it and so lose out on the money. Is it safe to apply the Willful Ignorance Principle in such cases?

The answer to this question is "it depends." What it depends upon is the impact of the belief that there is a duty not to inquire on D's t_1 psychology. In particular, sometimes such a belief will have a strong influence on the agent's attention. Someone who feels it is his duty not to inquire may put out of his mind the harms that might attend his action should the relevant condition be in place. He might put them out of his mind because he thinks he has no right to consider them in his calculations about what to do. In the case of the pharmacist, for instance, he might think that he would be violating his patients' right to privacy not just by inquiring but even by considering the question. Alternatively, he might put them out of his mind because he thinks that, given his duty not to inquire, he will not be responsible for the harms should they come to pass. In either of these cases, the Willful Ignorance Principle is not safely applied for the reasons it applies in core cases. Put in our terms, the reason is that it is critical to the application of the Willful Ignorance Principle in such cases that the t_1 a value be high; quite often, belief that it is one's duty not to inquire will undercut one's attention to the harms that attend t_2 action.

But, conversely, if the duty not to inquire does not drive down the a value, and if the omission of inquiry is indeed explicable by appeal to the way in which inquiry would place the benefit of action at risk, then the Willful Ignorance Principle is safely applied. Such cases differ from core cases in which there is no duty to omit inquiry, but not in a way that matters for our purposes here. There is, therefore, a class of cases in which the agent has a duty to omit inquiry, and fulfills that duty thus securing his ignorance, and yet is as culpable as a knowing actor.

5.2.2 False Positives and False Negatives

It is a defining feature of the core case that D believes that inquiry will succeed: he thinks that the probability that X is p and that the probability is p that inquiry will discover X. He

¹⁸ Deborah Hellman (2009), "Willfully Blind for Good Reason," *Criminal Law and Philosophy* 3(3): 301–316, makes the point that actors who omit inquiry in order to fulfill an obligation not to inquire are wrongly taken to be as culpable as the knowing. What is provided here can be taken as an explanation for why this is so.

thinks that the probability of either a false positive (where inquiry says that X when not- X) or a false negative (where inquiry says that not- X when X) is zero. Can we apply the Willful Ignorance Principle when this is not so? Say, for instance, that a friend asks D to sell him a case of pseudoephedrine that he happens to have in his storeroom. D is certain that, if he asks the friend if he plans to manufacture methamphetamine with the pseudoephedrine, the friend will say he does not. The question is whether D will be able to tell if the friend is lying. D might recognize, even if dimly, that there is both a non-zero false positive rate—sometimes he thinks people are lying who are telling the truth—and a non-zero false negative rate—sometimes he thinks people who are telling the truth are lying. He fails to inquire. The friend uses the drug to make methamphetamine and D is charged as an accomplice to the production on a willful ignorance theory. Is the Willful Ignorance Principle properly applied in such a case?

One class of agents will incorporate information about the false positive and false negative rates into their judgments of the probability that X is in place. That is, they will assess the probability of X to be just the same as the probability that their method of inquiry will say that X is in place. This is irrational for an agent who thinks the false positive and false negative rates to be non-zero; but especially when it comes to the assessment of probability, human beings are often irrational. For agents of this kind, the Willful Ignorance Principle is safely applied. Their cases are exactly like core cases in the relevant respect: they take the probability that X and the probability that inquiry will say that X to be the same.

The more interesting and harder cases are those in which D estimates the probability of X to be different from his estimate of the probability that inquiry will reveal X . This is an agent who makes estimates of the false positive and false negative rate; or rather, this is an agent whose estimates of those rates play a role in explaining his decision not to inquire. To see how things look for such an agent, consider a modified version of Table 1, which described the payoffs in the core case. Here, f_{neg} is the false negative rate (the probability that when X is in place inquiry will say not- X) and f_{pos} is the false positive rate (the probability that when not- X inquiry will say that X).

As before, we can use the values in this chart to calculate the ceiling on the agent's s value in light of the fact that he omits inquiry at t_1 . From the failure to inquire, together with the presumption of rationality, we infer the following:

$$\text{Value (not-inquire)} > \text{Value (inquire)}$$

Using the Value Equation, and the information in Table 2, we then reach the following conclusion:

$$(B * 1 - 1 * s * H * p) > (B * ((1 - p) - f_{pos} + f_{neg}) - 1 * s * H * f_{neg})$$

Some algebra leads to the following result:

$$s < B/H * ((p + f_{pos} - f_{neg}) / (p - f_{neg}))$$

Recall that the ceiling on the s value of the agent who acted knowingly at t_2 is B/H . So, we learn that the inferable ceiling on D 's s value is higher than that of the knowing agent just in case $((p + f_{pos} - f_{neg}) / (p - f_{neg})) > 1$, which is true just in case $f_{pos} > 0$.

What, exactly, does this result mean? To understand it, it helps to think about the meaning of false positives for an agent in a case that is otherwise like a core case. He thinks that if he believes that X is in place (e.g., that there are drugs in the car), then his conscience will not allow him to engage in the act that promises a benefit to him. He

Table 2 The values in this table are the same as in Table 1 with the exception of the four values associated with the conditional probabilities of arriving in a particular condition (e.g., X/Refrain), given the false positive (fpos) and false negative (fneg) rates

		X	Not-X
Act at t2	Outcome	B & H	B & not-H
	Prob Inquiry at t1	$0 + \text{fneg}$	$(1-p) - \text{fpos}$
	Prob No Inquiry	p	1-p
Refrain at t2	Outcome	not-B & not-H	not-B & not-H
	Prob Inquiry at t1	$p - \text{fneg}$	$0 + \text{fpos}$
	Prob No Inquiry	0	0

anticipates, then, that in the case of a false positive, he forgoes the benefit in an instance in which no harm would befall others were he to act. So, it is legitimate for him to take the possibility of a false positive to provide a reason against inquiry. Imagine, for instance, that inquiry were guaranteed to say that X was present, regardless of the facts. In that case, the false positive rate would be 1: the probability of inquiry saying that X, given not-X, would be 1. But if the agent were then to fail to inquire for that reason, he could not be faulted. Insofar as he is justified in pursuing the benefit when there is no prospect of harm for others, he is justified in eschewing inquiry for the reason that it might mislead him into avoiding fully justified conduct. This fact is accommodated in the model for the reasons described in the previous paragraph: the larger the false positive rate, the less fault there is in the agent's social preference manifested in his failure to inquire.

Notice, a parallel point cannot be made about the false negative rate. If the agent inquires and inquiry yields a false negative, then the agent will act and so will receive the benefit of action, but he will also inflict the harm on others. So, he should be attracted to inquiry that yields false negatives to the degree to which he is attracted to the benefit of action, and he should be averse to such inquiry to the degree to which he is averse to harming others. Since false negatives involve both the receipt of the benefit and the infliction of harm, the agent should be precisely as averse, all things considered, to a method of inquiry in light of its false negative rate as he is averse to engaging in the act knowing that X. This point, too, is incorporated into the model: the false negative rate, fneg, appears in *both* the numerator and the denominator in the term that modifies B/H. What this implies is that if the false positive rate (fpos) is zero, then the agent who omits inquiry is just as culpable as the knowing agent, even if the false negative rate is quite high. This is the right result.

In short, then, when the method of inquiry is imperfect, there are two important classes of cases: first, there are those in which D does not recognize inquiry's imperfection, but, instead, assigns the probability to X's occurrence to be the same as the probability that inquiry will say X. In such cases, the Willful Ignorance Principle is safely applied. Second, there are those cases in which D recognizes that there is a mismatch between the

probability that X and the probability that inquiry will say X. A subclass of these are cases in which inquiry's imperfection derives in part from its false positive rate, and that fact about inquiry is part of what explains why D omits it. In those cases, the Willful Ignorance Principle is not safe to apply; the failure to inquire just is not expressive of the same degree of moral failing that we find in the knowing agent. By contrast, in all other cases in that second class, including those in which the false negative rate is high, and that is part of the reason that D omits inquiry, the Willful Ignorance Principle is nonetheless safe to apply.

5.2.3 *Costly Inquiry*

Remember that in core cases the intrinsic costs of inquiry are negligible, or even zero. There is, to be sure, no shortage of cases in which this condition is not met; and if it is not, the case is not a core case. However, as I will now argue, a large class of these is nonetheless sufficiently similar to core cases to allow application of the Willful Ignorance Principle.

How should we think about willful ignorance when the agent would have to expend resources, or risk harm, in order to inquire at t1, and omits inquiry in part so as to avoid such costs? The simplest way to conceptualize the costs of inquiry is to supplement the Value Equation with a term, Q, representing the costs of inquiry that would not be incurred by refraining from inquiry:

$$\text{Value (inquiry)} = (B * (1 - p) - 1 * s * H * 0) - Q$$

As before, we can conclude from the fact that D chose not to inquire that

$$\text{Value (not-inquiry)} > \text{Value (inquiry)}$$

And this implies:

$$(B * 1 - 1 * s * H * p) > (B * (1 - p) - 1 * s * H * 0) - Q$$

A little algebra leads us to the conclusion that

$$s < B/H + Q/(H * p)$$

The final term in this equation is the crucial one. $Q/(H * p)$ can be thought of as a measure of the degree to which the fact that inquiry is costly counts in D's favor when he fails to inquire and acts at t2. It is a measure, that is, of what the costs of inquiry tell us about D's s level at t2.

Now, recall, that the s of an actor who knows that X at t2, when he acts, can be inferred to be below B/H . So, when inquiry is not costless, D's inferable s level is higher than that of the knowing actor. If we were incautious, we might conclude that when inquiry is costly there is something wrong with the Willful Ignorance Principle. After all, haven't we just shown that in such a case we have reason to believe that D cares more about others than the knowing actor? But this would be too quick. Rather, the right way to think about the issue is by considering the following question: say that D knows that X and A's at t2 and part of his reason for A-ing, rather than refraining from A, is that by A-ing he can avoid a result like that which would follow from inquiry, but with value $Q/(H * p)$. Would D, in that hypothetical, be diminished in his responsibility for knowingly A-ing?

For instance, imagine that it would cost D \$5 to inquire, that he thinks that the probability of X is .1, and that his A-ing will cost others \$1000, if X. The question, then, is this: would an agent who knowingly A'd (A'd while knowing X) be diminished in his

culpability had he done it, in part, to avoid paying $\$5/(\$1000 * .1) =$ five cents? The answer is “no” and so we learn that the agent who A’s at t2 without inquiring at t1 in such circumstances is properly treated as though he knowingly A’d at t2. That is, the Willful Ignorance Principle is properly applied in this case. Or, put another way, a case that differs from a core case only in that inquiry is costly differs in a way that matters for our purposes only when the costs of inquiry bear on his s level in a way that we are ready to credit him for. If we are not, then the Willful Ignorance Principle continues to be safe to use.

Consider another example: imagine that it would cost D \$5 to inquire, that he thinks the probability of X is .1, and that his A-ing will kill someone, if X. So, the measure of value that matters to us is this: $\$50/(\text{the value of a life})$. Although there are tools for measuring the value of a life in monetary terms, we needn’t employ them to think through the question of whether to treat D as knowing that he would kill someone when he A’d, at t2, in ignorance and, sadly, caused a death. Rather, we need to think through the culpability of a person who declines to pay \$50 to save a life. Does such a person have distorted values? Since the answer is “yes,” that’s a good reason to think that the willfully ignorant agent should be treated as though he were knowing in this case. Our theory does not obviate the need to make judgments about cases that seem to call for the aggregation of things of incommensurable value. Rather, the theory exploits our capacity to make such judgments, which we do all the time. If there are cases in which the call is harder—imagine, for instance, that one would have to spend one’s life savings to save a life—then the theory has the implication that the correlative implication about the applicability of the Willful Ignorance Principle will be just as hard.

Return for a moment to Heredia—the woman whose mother and aunt had her drive them across the border in the detergent-smelling car with a trunk full of marijuana. Why didn’t Heredia inquire further, before it was too late, about the source of the smell and the source of her mother’s surprisingly large stack of cash? It is hard to know for certain; she may not know herself; but one possibility is that she feared what such an inquiry would do to her relationship with her mother and her aunt. If she settled the question and discovered that there were indeed drugs in her trunk, then she would need to either embrace the criminal scheme, or else, we can imagine, precipitate a break with her relatives. So, if this is the right reconstruction of her psychology, what she was unwilling to do was to place the benefits of her family associations at risk while recognizing that without inquiry that same degree of risk would attend the possibility that she would move drugs across the border.

So, for Heredia, inquiry was not costless. The inquiry itself would harm her relationship with her mother and her aunt, even if it turned out that there were no drugs in the car. We can use the analysis just described to determine whether it is appropriate to treat her as if she knew that there were drugs in the car when she drove across the border. The relevant quantity is:

$$\begin{aligned} & (\text{the value of the harm to her familial relationships}) / \\ & ((\text{the harm involved in transporting the drugs}) \\ & * (\text{the probability that there were drugs in the car})) \end{aligned}$$

The question is whether someone who knew that there were drugs in the car, and drove them across the border, in part so as to avoid incurring a cost of this magnitude, would be thereby diminished in culpability for bringing the drugs across the border, or would be, instead, just as culpable as an agent who knew the drugs were in the car and was trying instead solely to achieve the goods of the payoff of the action to himself. And relevant to this question is what we would think of the values of someone willing to risk bringing the

drugs into the country so as to avoid harm to her familial relationships. What we would think of such a person tells us something about how we think people ought to weigh the values in the numerator and the denominator of the relevant quantity. And that in turn tells us whether someone who weighed them in such a way as to act for the sake of avoiding a cost of the relevant magnitude is less culpable than someone who does not consider such a cost in deciding to act.

Notice, also, that the probability that there are drugs in the car figures in the denominator of the relevant quantity. This is important, for if the probability is low enough, the result is that this quantity is enormous. That is, if the probability is low enough, what we are to ask ourselves is whether someone who knowingly moved drugs across the border in order to avoid paying *an enormous cost* is reduced in culpability in comparison to someone who knowingly moved the drugs across the border without such a prospect. If the probability is low enough, that is, the answer to this question very well might be “yes.” And, if it is, then the Willful Ignorance Principle does not apply; the agent is not equal in culpability to the knowing actor. Heredia’s is a sympathetic case, I believe, because it is plausible to believe both that her familial relationships were of the utmost importance, and that her estimates of the probability that her car contained drugs were low. Add, what many of us feel, that there is little real harm in transporting marijuana into the United States, and it can seem irresistible to conclude that she is less culpable than the knowing actor. The costs of inquiry were high enough that even a knowing actor who moved drugs across the border to avoid them would have been reduced in culpability. Perhaps. There are, however, several empirical assumptions here, and whether they are true of Heredia is hard to discern.

The approach just offered for understanding the bearing of a failure to engage in costly inquiry to willful ignorance involves a normative assumption. The assumption is that we ought to separate, in our thinking about a person’s responsibility for wrongful behavior, those goods, on the one hand, that he pursued *through* the act (what we have been calling EV_{Self}) and, on the other, those costs that he would incur were he to refrain from the act (here, Q , or the costs of inquiry). The claim is that we make an initial inquiry into D ’s s value leaving out the role of Q in driving his behavior and then ask whether it in any sense ameliorates his responsibility that part of the explanation for the fact that he omitted inquiry was that he thereby avoided paying $Q/(H*p)$. Put in legal terms, the outcome of applying the Value Equation without consideration of the costs of inquiry supports something like a *prima facie* case for treating him as though he were knowing, while consideration of the costs of inquiry, if it counts in his favor, does so in something like the way that an affirmative defense does. If (a) he omitted inquiry and later acted in ignorance and, (b) his would be a core case but for the fact that inquiry was not costless, then he must now explain why he is not to be treated as though he were knowing. He can provide such an explanation by showing that those who knowingly act to avoid a harm of the relevant magnitude ($Q/(H*p)$) are diminished in culpability in a way that warrants treating them differently from other knowing actors. And when he cannot show that, then he is properly treated as though he were knowing. This way of conceptualizing the matter, then, reveals a commitment of the present approach: there is a moral fact of the matter about which factors are appropriately considered to be part of what people ought to answer for, on the one hand, and what factors are appropriately considered to be part of the answer they might give in their defense. This is not just a matter of legislative drafting; it is not a purely discretionary choice, but, instead, can be done well or badly.¹⁹ Or, at least, that is a presupposition of the present approach.

¹⁹ The point here strongly echoes ideas presented in R. A. Duff (2007), *Answering for Crime: Responsibility and Liability in the Criminal Law*, Oxford: Hart Publishing.

5.2.4 Foxes

The approach just offered for conceptualizing cases that would be core cases but for the fact that inquiry is costly can help us to understand what is going on in cases of defendants who omit inquiry *so as to avoid prosecution later*. Some commentators have taken these to be the core cases of willful ignorance—the cases that they take the Willful Ignorance Principle to be designed for. These are the agents David Luban has called the “foxes.”²⁰ If anyone should be treated as if they were knowing, the foxes should; this is a fact that any rationale for the Willful Ignorance Principle should accommodate. Now, notice that the paradigm fox does not anticipate that he will fail to perform the act if he has the knowledge that inquiry would bring. The paradigm fox is planning to perform the act *regardless* of whether he knows that X. He just believes (falsely, given that the Willful Ignorance Principle is law) that A promises greater benefit if he does not inquire than if he does. If he inquires, he thinks, then A promises its benefits minus a risk of punishment for a knowing crime, while if he does not inquire, then A promises those benefits without the same accompanying risk. So, it is very easy to explain why the Willful Ignorance Principle applies to the *paradigm* fox: the fact that one anticipates a risk of apprehension and punishment when one performs the act knowingly shows that the EV_{Self} term in the Value Equation is lower than if one did not anticipate that. And so, in such a case, the agent’s s value has to be even lower in order to explain why the act was worth doing. Anticipation of the possibility of sanction that fails to deter, that is, inculcates rather than exculpates. And since we know that the paradigm fox will happily perform the act with knowledge, the fact that he tries by omitting inquiry to make the act more valuable to him does not count in any way in his favor. It does not show his s value to be any higher than the knowing actor.

However, not every fox is a paradigm fox. Some subset of the foxes should be conceptualized as just like core cases in which inquiry is not costless. Imagine someone who omits inquiry out of fear of a risk of punishment for the knowing crime *and would not perform A at t_2 were he to know that X*. His would be an example of the core case but for the fact that inquiry comes with a cost: it removes the possibility, he thinks, of A-ing at t_2 free of the possibility of punishment for knowingly A-ing. Using the foregoing analysis, the question of whether to apply the Willful Ignorance Principle reduces to this: imagine that D were to commit the knowing crime so as to avoid a cost of the following magnitude: (the cost of the risk of apprehension and punishment for the knowing crime)/(the expected harm of the criminal act). Would an agent have reduced culpability for criminal behavior performed to avoid a cost of that size? We can begin to address this question, in turn, by contemplating the agent who imposes the expected harm of the criminal act on others in order to avoid an expected punishment to himself. Since, on the assumption that punishments are proportional to the crimes for which they are issued, such a person clearly has distorted values, we should conclude that it is safe to employ the Willful Ignorance Principle even with respect to non-paradigmatic foxes whose cases would be core cases but for the fact that inquiry is costly.

6 Conclusion

Under the rationale offered here for the Willful Ignorance Principle, *mens rea* matters to the criminal law’s response to the occurrence of a complex set of conditions that includes it because in light of our mental states our conduct manifests something more fundamental to

²⁰ David Luban (1999), “Contrived Ignorance,” *Georgetown Law Review* 87: 957–980.

culpability: our social preferences. When we give this idea expression in the Value Equation, we find that it is not just plausible, but *demonstrable* that the person who sees the receipt of the benefit of action as placed at risk by inquiry, and so does not inquire, manifests precisely the same social preference as the person who flat out knows that his conduct will give rise to the violation of the relevant legally protected interest when he acts. This should give us confidence about our use of the Willful Ignorance Principle—at least when we apply it in core cases.

One argument for a theory of the point of *mens rea*, such as the theory offered here, is that the theory provides an explanation for something that was pre-theoretically appealing, but hard to explicitly justify. That, in fact, is the argument offered here for the theory of *mens rea*'s point offered in Sect. 3. The fact that that theory provides us with a simple and illuminating rationale for the Willful Ignorance Principle counts in favor of the theory, given that there simply must be a domain in which it is appropriate to apply the Willful Ignorance Principle.

But this is not the only reason to be attracted to the theory of *mens rea*'s point offered here. Criminal law exhibits a well-known fixation with the defendant's mind, a fixation that we do not find in other areas of law, including areas in which the mental states of the parties matter to liability. Criminal law responds differently to defendants who are only subtly different in their psychological states; we often give large punishments to some who cause harm while giving low punishments, or even no punishments, to subtly psychologically different actors who cause the very same, or even greater, harms. The reason is that subtle differences in psychological states—such as that between knowing that X is in place and merely recognizing a very good chance of that—cloak large differences in something fundamental to what it is to be a human being and a citizen of a state who owes an account of his conduct to other people and other citizens: the evaluative weight that we give to others' interests in comparison to our own. They cloak facts about the ways in which people employ their rational and evaluative capacities, the very capacities thanks to which they are capable of appreciating the reasons for the demands to which they are subject. From subtle differences in psychology, we are able to infer the presence of large differences in the facts about how an agent employs the very capacities that make him a human agent, worthy of punishment and censure for wrongdoing. Why else would we care about small differences in psychological outlook? What else could the point of *mens rea* be?²¹

²¹ This material has been presented in several places and I received very valuable comments on each occasion. Audiences at Fordham, Georgia State University, University of California at Irvine, University of Connecticut, University of Michigan, and Yale were all very helpful. Thanks are particularly owed to Anne Dailey, Debby Denno, Sharon Dolovich, Bill Edmonson, Jeff Helmreich, Issa Kohler-Hausmann, David Manley, Tracey Meares, Gabe Mendlow, Eddy Nahmias, Jacob Ross, Alex Sarch, and Alan Schwartz.